

K-Nearest Neighbors (KNN) Classification Project

Author: Ayoub Majjid

Email: ayoub@majjid.com

Portfolio: majjid.com

Repository: <https://github.com/ayoubmajid67/learn.git>

Project Goal

This project predicts **whether a user will purchase a product** based on demographic data such as:

***Age**

***Estimated Salary**

***Gender**

We use the **K-Nearest Neighbors (KNN)** algorithm — a simple but powerful **machine learning model** that classifies a new point based on its **closest neighbors** in the training data.

Notebook Outline

This notebook demonstrates how to:

1. Load and explore data
 2. Preprocess data for ML models
 3. Train the KNN classifier
 4. Evaluate model performance
 5. Visualize decision boundaries
 6. Interpret the results visually and statistically
-

Step 0: Understanding the Idea Behind the KNN Algorithm

Before we start coding, let's understand **what KNN is and how it works**.

What is KNN?

K-Nearest Neighbors (KNN) is one of the simplest **supervised machine learning algorithms** used for **classification** and **regression**.

- “Supervised” means it **learns from labeled examples** — data where we already know the correct answer (for example, whether a user bought a product or not).
- KNN doesn’t build an equation or internal model. Instead, it **stores all training examples** and classifies new data by **comparing** it to those examples.

It's often described as "**learning by analogy**."

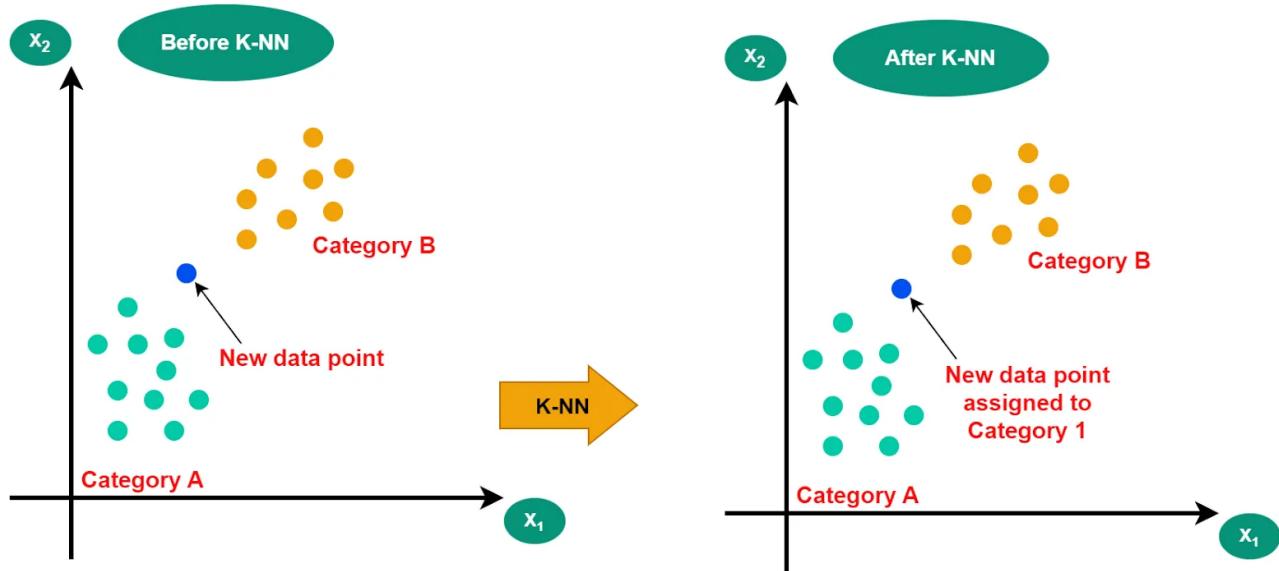
💡 How It Works – Step by Step

Imagine you have a dataset of users with:

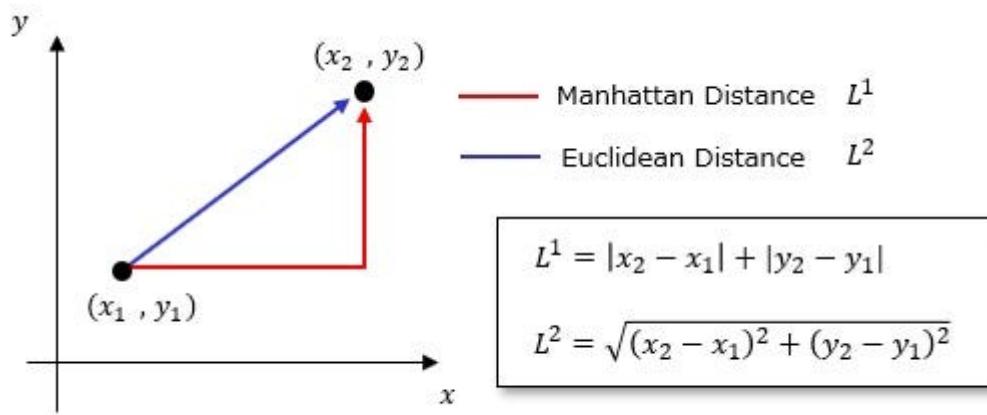
- Age
- Estimated Salary
- Whether they purchased a product (Yes/No)

Now, suppose we meet a **new user** and want to predict if they will buy the product.

Here's what KNN does:



1 Compute Distances:



1 Compute Distances:

- Measure how "close" the new user is to every user in the training set.
- The most common distance metric is the **Euclidean distance**:

[

$$d = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2 + \dots}$$

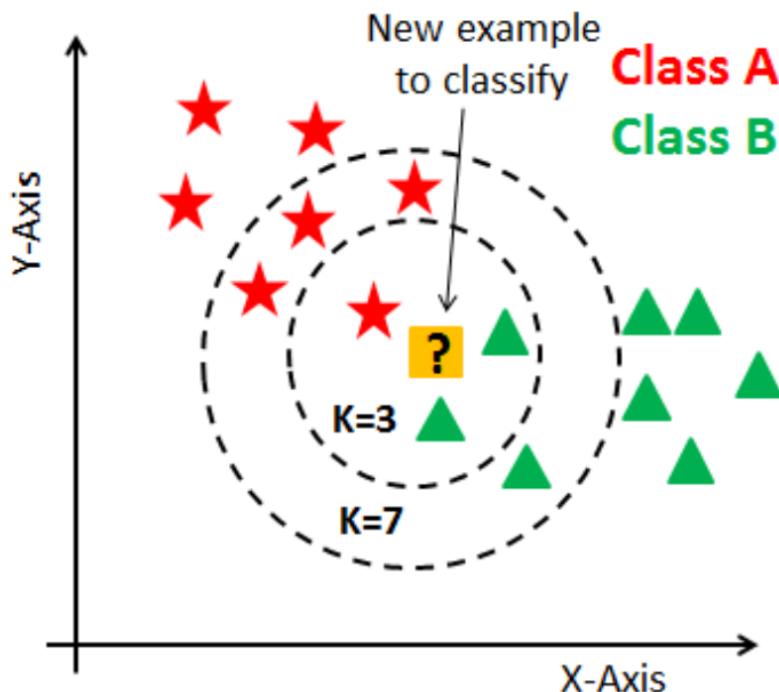
]

This tells us how far points are in multi-dimensional space.

2 Find the K Nearest Neighbors:

- Sort all training samples by distance from the new user.
 - Pick the **K closest ones** (for example, the 5 nearest points).
-

3 Voting:



- Each neighbor "votes" for its class (Purchased = 1, Not Purchased = 0).
- The class with the **most votes** becomes the prediction.

Example:

$$K = 5 \rightarrow [1, 0, 1, 1, 0]$$

Votes \rightarrow Purchased = 3, Not Purchased = 2 \rightarrow Predict Purchased (1)

⚙️ Choosing the Right Value of K

***Small K** → very sensitive to noise (too local, may overfit).

***Large K** → smoother, but may ignore local patterns (underfit).

- Typically, we test several values of K and pick the one with the best validation accuracy.
-

⌚ Why Does KNN Work Well?

KNN works best when:

- The **data has clear clusters or groups**, like "buyers" and "non-buyers."
- The **features are scaled**, so no feature dominates distance calculation.
- The **relationship between inputs and output is non-linear** — KNN doesn't assume any mathematical formula, it adapts to the data shape.

In our case:

- Users with **similar ages and incomes** tend to behave similarly.
 - So KNN can easily find patterns based on proximity in feature space.
-

❖ KNN in One Sentence:

"To predict something new, look at your closest examples and do what most of them do."

Now that you understand the concept, let's move on to **data loading and preprocessing** so we can apply this algorithm in practice.

☒ Step 0.5: Understanding the Project Structure

```
KNN_Project/
|
└── data/
    └── Social_Network_Ads.csv      # Dataset used for training and testing
|
└── docs/
    ├── api_documentation.pdf      # API documentation (PDF)
    └── project_overview.md        # General project explanation
|
└── KNN_SIG_MODEL.postman_collection.json # Postman collection for testing
```

API

```
|  
|   └── models/  
|       └── knn_model.pkl          # Saved trained KNN model for reuse  
|  
|  
|   └── notebooks/  
|       └── knn_classification.ipynb    # Main Jupyter Notebook  
|  
|  
└── src/                          # Source code directory containing  
    reusable modules  
    |   ├── __init__.py            # Marks src as a Python package  
    |   ├── data_loader.py        # Loads and returns the dataset  
    |   ├── preprocessing.py     # Encodes, scales, and splits the data  
    |   ├── model_train.py       # Trains and saves the KNN model  
    |   ├── evaluate.py          # Evaluates the model (accuracy,  
    |   |   confusion matrix)  
    |   └── visualize.py        # Plots decision boundaries and data  
        distributions  
    |  
    |  
    └── api/                      # Flask API application  
        |   ├── __init__.py          # Initialize the Flask app  
        |   └── routes/              # Endpoints of the API  
        |       └── api_bp.py        # Blueprint with all endpoints  
    |  
    |  
    └── main_train.py            # Script to train the model outside  
        notebook  
    └── main.py                  # Start the Flask API server  
    └── requirements.txt         # Dependencies list (e.g., pandas,  
        sklearn, flask)
```

```
└─ README.md # Project overview, usage  
    instructions, setup guide
```

🔍 Explanation of Each Component

-data/ → Stores all raw or processed datasets used by the model.

-models/ → Keeps trained model files so they can be reused in APIs or other applications.

-src/ → Contains all modularized Python scripts for each process: loading, preprocessing, training, evaluation, and visualization.

-notebooks/ → Interactive notebooks used for experimentation and documentation.

-requirements.txt → Lists all Python dependencies to replicate the environment easily.

This modular design follows the principle of **separation of concerns**, ensuring that each part of the system has a single responsibility.

Step-by-Step Guide to Run the Project (Using **venv**)

1. Clone the repository:

```
git clone https://github.com/ayoubmajid67/learn.git
```

```
cd projects/KNN_Project
```

2. Create a virtual environment:

```
python -m venv venv
```

3. Activate the virtual environment:

*Windows (cmd):

```
venv\Scripts\activate
```

***Windows (PowerShell):**

```
venv\Scripts\Activate.ps1
```

***Linux / MacOS:**

```
source venv/bin/activate
```

4. Install dependencies:

```
pip install -r requirements.txt
```

5. Train the model (optional if not already trained):

```
python main_train.py
```

- Reads the dataset, preprocesses it, trains the KNN model, and saves it to `models/knn_model.pkl`.

6. Start the Flask API server:

```
python main.py
```

- The API will run at `http://127.0.0.1:5000`.

7. Test the API:

***Check API documentation:** Open `docs/api_documentation.pdf`.

***Postman testing:** Import `docs/KNN_SIG_MODEL.postman_collection.json` into Postman and run the endpoints directly.

- Example request :

The screenshot shows the Postman interface. At the top, it says "HTTP KNN_SIG_MODEL / predict". Below that, a "POST" request is selected with the URL "localhost:5000/api/predict". The "Body" tab is open, showing a JSON payload with fields "Age" (30), "EstimatedSalary" (50000), and "Gender" (1). The response tab shows a status of 200, with the message: "The model predicts the user will NOT purchase.", and a prediction value of 0.

Tips / Notes

- Keep your virtual environment **activated** whenever you work on the project.
- Store **all new documentation or API updates** in the `docs/` folder.
- The **Postman collection** is a ready-to-use way to test all endpoints without writing additional code.

Step 1: Import Dependencies and Setup

We'll import the modules from our `/src` package to keep our project modular and clean.

```
import sys
import os

# Optional: run once to include project root in sys.path
"""

project_root = os.path.abspath(os.path.join(os.getcwd(), '..'))

if project_root not in sys.path:

    sys.path.append(project_root)
"""

from src.model_train import train_knn, save_model
```

```
from src.evaluate import evaluate_model

from src.data_loader import load_data

from src.preprocessing import preprocess_data

from src.visualize import plot_decision_boundary_train_test

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns
```

📊 Step 2: Load and Explore the Dataset

We start by loading our dataset — **Social_Network_Ads.csv**, which contains information about users and whether they purchased a product after seeing an ad.

```
def load_data(path='data/Social_Network_Ads.csv'):

    return pd.read_csv(path)

df = load_data("../data/Social_Network_Ads.csv")

df.head()

df.info()
```

⌚ Observations

***Features:** Age, Gender, EstimatedSalary

***Target:** Purchased (1 = yes, 0 = no)

- We will later convert categorical values like "Male/Female" into numeric codes.

🛠️ Step 3: Data Preprocessing

Machine Learning models can't understand text or unscaled values — they work best with **numerical, normalized data**.

```

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder, StandardScaler


def preprocess_data(df):

    # Encode gender (text → numeric)

    df['Gender'] = LabelEncoder().fit_transform(df['Gender'].astype(str))

    # Split into features and target

    X = df[['Age', 'EstimatedSalary', 'Gender']].values

    y = df['Purchased'].values

    # Train-test split (75% training, 25% testing)

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=0)

    # Scale features for KNN (important!)

    sc = StandardScaler()

    X_train = sc.fit_transform(X_train)

    X_test = sc.transform(X_test)

    return X_train, X_test, y_train, y_test

X_train, X_test, y_train, y_test = preprocess_data(df)

print(f"Training set: {X_train.shape}, Test set: {X_test.shape}")

```

💡 Why Scaling?

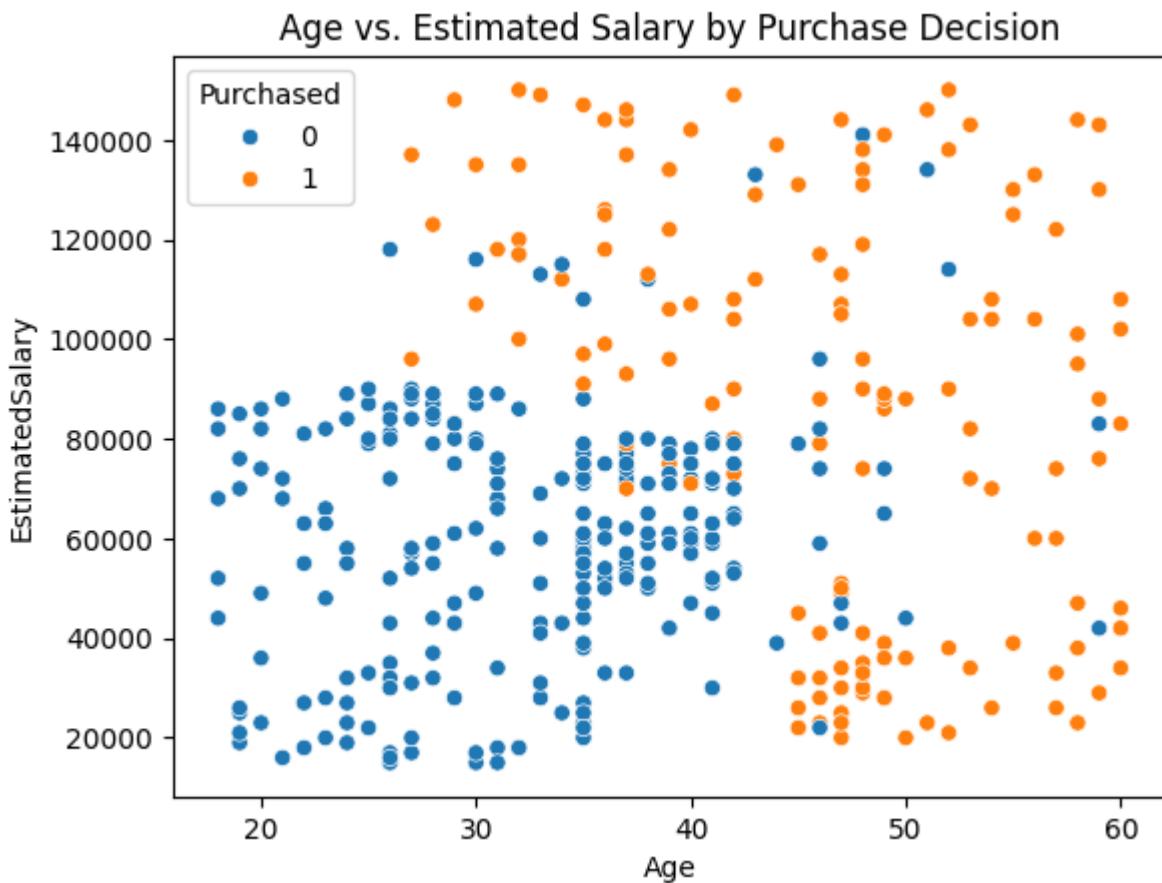
KNN uses **distance** to decide neighbors.

Without scaling, a feature like *salary* (which can be in thousands) would dominate over *age*.

– Step 4: Visualize the Data

Before training, let's visualize how Age and Estimated Salary relate to purchasing decisions.

```
sns.scatterplot(data=df, x='Age', y='EstimatedSalary', hue='Purchased',  
palette='coolwarm')  
  
plt.title("Age vs. Estimated Salary by Purchase Decision")  
  
plt.show()
```



🧠 Interpretation:

- Clusters form naturally: younger users with lower income rarely purchase.
- Older or higher-salary users are more likely to purchase.
- This gives intuition for why KNN (a distance-based algorithm) works well here.

🍪 Step 5: Train the KNN Model

KNN works by looking at the **K nearest neighbors** of a point and letting them "vote" on its class.

```
model = train_knn(X_train, y_train)
```

⌚ Example:

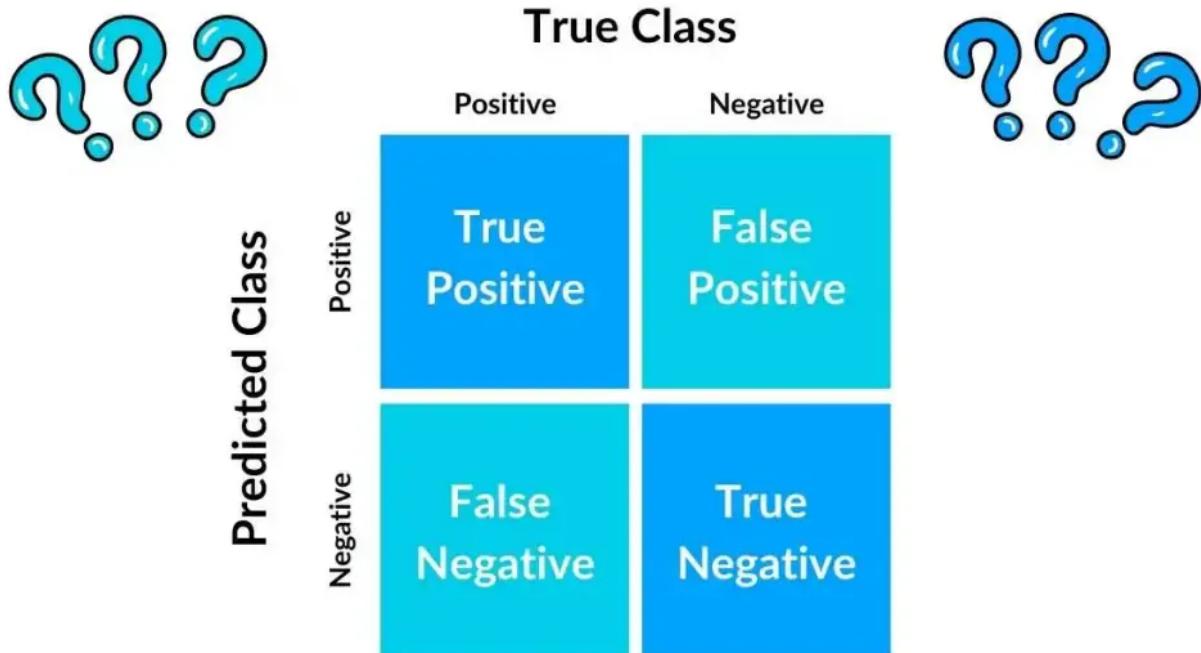
If (K = 5) and among the 5 nearest users, 3 purchased and 2 did not → the model predicts **Purchased = 1**.

💾 Step 6: Save the Trained Model

We save the trained KNN model using Python's **pickle** for reuse in APIs or apps.

```
save_model(model, '../models/knn_model.pkl')
```

📝 Step 7: Evaluate the Model



We test our model on unseen data and calculate performance metrics.

```
acc, cm = evaluate_model(model, X_test, y_test)

print(f"Accuracy: {acc:.2f}")

print("Confusion Matrix:\n", cm)
```

```

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

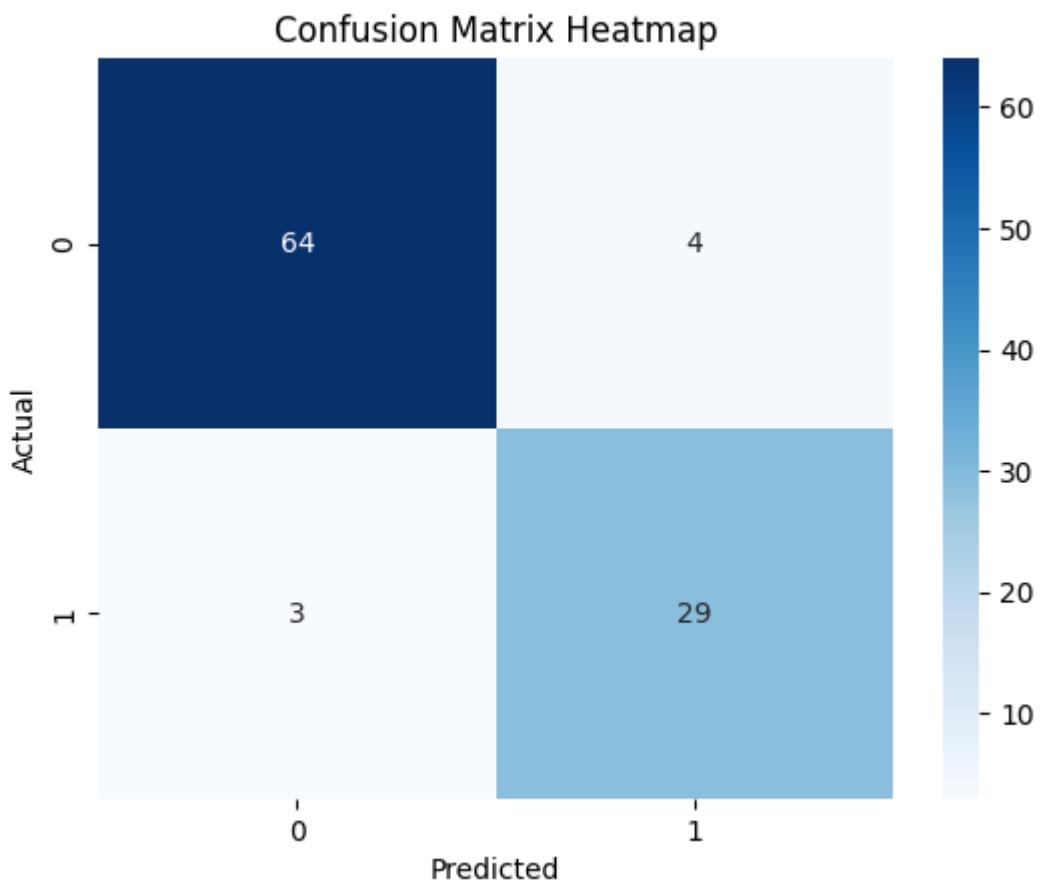
plt.title("Confusion Matrix Heatmap")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()

```



Interpreting the Confusion Matrix

From your documentation:

```

[[64  4]
 [ 3 29]]

```

| Term | Meaning | Count |

| ----- | ----- | ----- |

| **TN (True Negative)** | Correctly predicted non-buyers | 64 |

| **FP (False Positive)** | Incorrectly predicted buyers | 4 |

| **FN (False Negative)** | Missed actual buyers | 3 |

| **TP (True Positive)** | Correctly predicted buyers | 29 |

Accuracy: 93%

Precision: 88%

Recall: 91%

F1-Score: 89%

The model performs well — minimal misclassifications, strong generalization.

⌚ Step 8: Visualize Decision Boundaries

Now, let's see **how the model makes its decisions visually**.

```
plot_decision_boundary_train_test(model, X_train, y_train, X_test, y_test)
```

⌚ How It Works (from your visualization doc):

1. Meshgrid Creation – a grid of possible (Age, Salary) pairs is generated.

2. Model Prediction – KNN predicts each point as "buy" or "not buy."

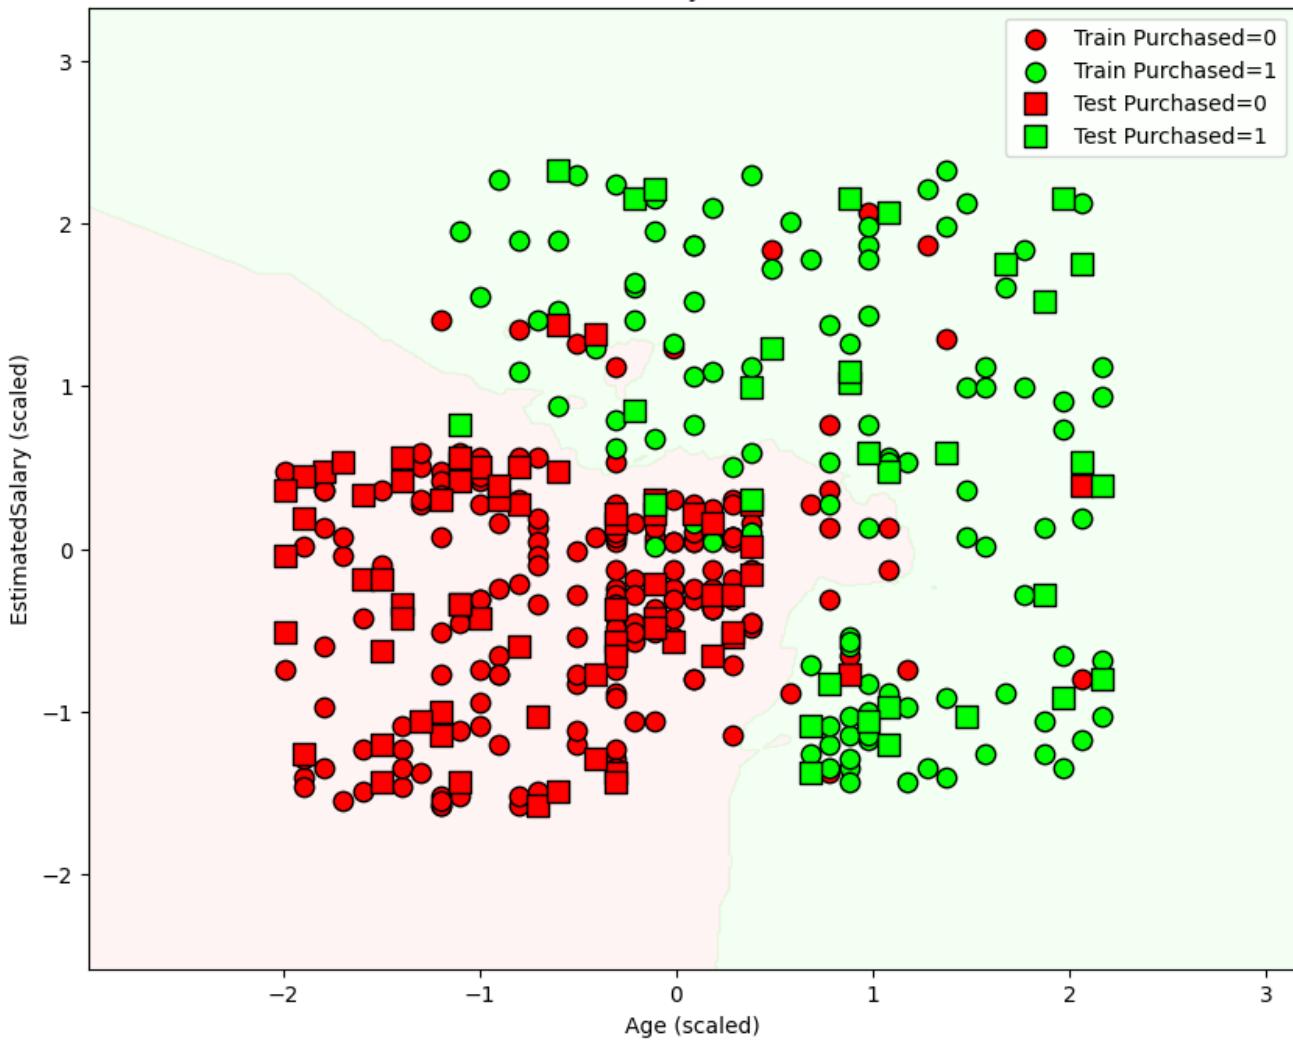
3. Color Regions:

-  **Red:** Model predicts "Not Purchased (0)"
-  **Green:** Model predicts "Purchased (1)"

4. Data Points:

- Circles → Training set
- Squares → Test set

KNN Decision Boundary with Train & Test Sets



⌚ Graph Insights

- Green and red zones show **decision boundaries** created by KNN.
- Overlapping regions represent **uncertain areas**.
- Misclassified points appear in the “wrong” color zone.
- The model captures complex shapes — a major strength of KNN.

✓ Step 9: Conclusion

***KNN** effectively distinguishes buyers from non-buyers based on demographic data.

***Accuracy:** 93% with balanced precision and recall.

***Visualization** reveals how KNN adapts to nonlinear patterns.

***Model saved** as `models/knn_model.pkl` for reuse.

🌐 Author Information

 **Author:** Ayoub Majjid

 **Email:** ayoub@majjid.com

 **Portfolio:** majjid.com

** Repository: <https://github.com/ayoubmajid67/learn.git>