

K-Nearest Neighbors (KNN) Classification Project

Author: Ayoub Majjid

Email: ayoub@majjid.com

Portfolio: majjid.com

Repository: <https://github.com/ayoubmajjid67/learn.git>

Project Goal

This project predicts **whether a user will purchase a product** based on demographic data such as:

*Age

*Estimated Salary

*Gender

We use the **K-Nearest Neighbors (KNN)** algorithm — a simple but powerful **machine learning model** that classifies a new point based on its **closest neighbors** in the training data.

Notebook Outline

This notebook demonstrates how to:

1. Load and explore data
 2. Preprocess data for ML models
 3. Train the KNN classifier
 4. Evaluate model performance
 5. Visualize decision boundaries
 6. Interpret the results visually and statistically
-

Step 0: Understanding the Idea Behind the KNN Algorithm

Before we start coding, let's understand **what KNN is and how it works**.

What is KNN?

K-Nearest Neighbors (KNN) is one of the simplest **supervised machine learning algorithms** used for **classification** and **regression**.

- “Supervised” means it **learns from labeled examples** — data where we already know the correct answer (for example, whether a user bought a product or not).
 - KNN doesn't build an equation or internal model. Instead, it **stores all training examples** and classifies new data by **comparing** it to those examples.
-

It's often described as “**learning by analogy.**”


How It Works – Step by Step

Imagine you have a dataset of users with:


- Age
- Estimated Salary
- Whether they purchased a product (Yes/No)

Now, suppose we meet a **new user** and want to predict if they will buy the product.

Here's what KNN does:

 1760812080110

1 Compute Distances:

 1760812104391

1 Compute Distances:

- Measure how “close” the new user is to every user in the training set.
- The most common distance metric is the **Euclidean distance**:

[

$$d = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2 + \dots}$$


]

This tells us how far points are in multi-dimensional space.

2 Find the K Nearest Neighbors:

- Sort all training samples by distance from the new user.
 - Pick the **K closest ones** (for example, the 5 nearest points).
-

3 Voting:

 1760812117847

- Each neighbor “votes” for its class (Purchased = 1, Not Purchased = 0).
- The class with the **most votes** becomes the prediction.

Example:

$$K = 5 \rightarrow [1, 0, 1, 1, 0]$$

Votes → Purchased = 3, Not Purchased = 2 → Predict Purchased (1)

⚙️ Choosing the Right Value of K

***Small K** → very sensitive to noise (too local, may overfit).

***Large K** → smoother, but may ignore local patterns (underfit).

- Typically, we test several values of K and pick the one with the best validation accuracy.

🔗 Why Does KNN Work Well?

KNN works best when:

- The **data has clear clusters or groups**, like “buyers” and “non-buyers.”
- The **features are scaled**, so no feature dominates distance calculation.
- The **relationship between inputs and output is non-linear** — KNN doesn’t assume any mathematical formula, it adapts to the data shape.

In our case:

- Users with **similar ages and incomes** tend to behave similarly.
- So KNN can easily find patterns based on proximity in feature space.

🔗 KNN in One Sentence:

“To predict something new, look at your closest examples and do what most of them do.”

Now that you understand the concept, let’s move on to **data loading and preprocessing** so we can apply this algorithm in practice.

📁 Step 0.5: Understanding the Project Structure

```
KNN_Project/  
|  
├── data/  
|   └── Social_Network_Ads.csv           # Dataset used for training and testing  
|
```

```

├─ docs/                                # Documentation files
|   └─ api_documentation.pdf            # API documentation (PDF)
|   └─ project_overview.md              # General project explanation
|   └─ KNN_SIG_MODEL.postman_collection.json # Postman collection for testing API
|
├─ models/
|   └─ knn_model.pkl                    # Saved trained KNN model for reuse
|
├─ notebooks/
|   └─ knn_classification.ipynb         # Main Jupyter Notebook
|
├─ src/                                # Source code directory containing reusable modules
|   └─ __init__.py                      # Marks src as a Python package
|   └─ data_loader.py                  # Loads and returns the dataset
|   └─ preprocessing.py                # Encodes, scales, and splits the data
|   └─ model_train.py                  # Trains and saves the KNN model
|   └─ evaluate.py                     # Evaluates the model (accuracy, confusion matrix)
|   └─ visualize.py                    # Plots decision boundaries and data distributions
|
├─ api/                                # Flask API application
|   └─ __init__.py                     # Initialize the Flask app
|   └─ routes/                          # Endpoints of the API
|       └─ api_bp.py                   # Blueprint with all endpoints
|
└─ main_train.py                       # Script to train the model outside

```

notebook

├─ main.py # Start the Flask API server

├─ requirements.txt # Dependencies list (e.g., pandas, sklearn, flask)

└─ README.md # Project overview, usage instructions, setup guide

🔍 Explanation of Each Component

-**data/** → Stores all raw or processed datasets used by the model.

-**models/** → Keeps trained model files so they can be reused in APIs or other applications.

-**src/** → Contains all modularized Python scripts for each process: loading, preprocessing, training, evaluation, and visualization.

-**notebooks/** → Interactive notebooks used for experimentation and documentation.

-**requirements.txt** → Lists all Python dependencies to replicate the environment easily.

This modular design follows the principle of **separation of concerns**, ensuring that each part of the system has a single responsibility.

Step-by-Step Guide to Run the Project (Using **venv**)

1.Clone the repository:

```
git clone https://github.com/ayoubmajid67/learn.git
```

```
cd projects/KNN_Project
```

2.Create a virtual environment:

```
python -m venv venv
```

3.Activate the virtual environment:

*Windows (cmd):

```
venv\Scripts\activate
```

***Windows (PowerShell):**

```
venv\Scripts\Activate.ps1
```

***Linux / MacOS:**

```
sourcevenv/bin/activate
```

4.Install dependencies:

```
pipinstall-rrequirements.txt
```

5.Train the model (optional if not already trained):

```
pythonmain_train.py
```

- Reads the dataset, preprocesses it, trains the KNN model, and saves it to `models/knn_model.pkl`.

6.Start the Flask API server:

```
pythonmain.py
```

- The API will run at `http://127.0.0.1:5000`.

7.Test the API:

***Check API documentation:** Open [docs/api_documentation.pdf](#).

***Postman testing:** Import [docs/KNN_SIG_MODEL.postman_collection.json](#) into Postman and run the endpoints directly.

- Example request using `curl`:

```
curl-XPOSThttp://127.0.0.1:5000/predict\  
  
-H "Content-Type: application/json" \  
  
-d'{"Age": 35, "EstimatedSalary": 50000}'
```

☑ Tips / Notes

- Keep your virtual environment **activated** whenever you work on the project.
- Store **all new documentation or API updates** in the [docs/](#) folder.
- The **Postman collection** is a ready-to-use way to test all endpoints without writing additional code.

⚙ Step 1: Import Dependencies and Setup

We'll import the modules from our [/src](#) package to keep our project modular and clean.

```
import sys  
  
import os  
  
# Optional: run once to include project root in sys.path  
"""  
  
project_root = os.path.abspath(os.path.join(os.getcwd(), '..'))  
  
if project_root not in sys.path:  
    sys.path.append(project_root)  
"""  
  
from src.model_train import train_knn, save_model  
  
from src.evaluate import evaluate_model
```

```
from src.data_loader import load_data

from src.preprocessing import preprocess_data

from src.visualize import plot_decision_boundary_train_test


import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns
```

Step 2: Load and Explore the Dataset

We start by loading our dataset — **Social_Network_Ads.csv**, which contains information about users and whether they purchased a product after seeing an ad.

```
def load_data(path='data/Social_Network_Ads.csv'):

    return pd.read_csv(path)


df = load_data("../data/Social_Network_Ads.csv")

df.head()

df.info()
```

Observations

***Features:** Age, Gender, EstimatedSalary

***Target:** Purchased (1 = yes, 0 = no)

- We will later convert categorical values like “Male/Female” into numeric codes.

Step 3: Data Preprocessing

Machine Learning models can't understand text or unscaled values — they work best with **numerical, normalized data**.


```

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder, StandardScaler

def preprocess_data(df):

    # Encode gender (text → numeric)

    df['Gender'] = LabelEncoder().fit_transform(df['Gender'].astype(str))

    # Split into features and target

    X = df[['Age', 'EstimatedSalary', 'Gender']].values

    y = df['Purchased'].values

    # Train-test split (75% training, 25% testing)

    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
random_state=0)

    # Scale features for KNN (important!)

    sc = StandardScaler()

    X_train = sc.fit_transform(X_train)

    X_test = sc.transform(X_test)

    return X_train, X_test, y_train, y_test

X_train, X_test, y_train, y_test = preprocess_data(df)

print(f"Training set: {X_train.shape}, Test set: {X_test.shape}")

```

💡 Why Scaling?


KNN uses **distance** to decide neighbors.

Without scaling, a feature like *salary* (which can be in thousands) would dominate over *age*.

📊 Step 4: Visualize the Data

Before training, let's visualize how Age and Estimated Salary relate to purchasing decisions.

```
sns.scatterplot(data=df, x='Age', y='EstimatedSalary', hue='Purchased',  
palette='coolwarm')  
  
plt.title("Age vs. Estimated Salary by Purchase Decision")  
  
plt.show()
```

 1760813401423

Interpretation:

- Clusters form naturally: younger users with lower income rarely purchase.
- Older or higher-salary users are more likely to purchase.
- This gives intuition for why KNN (a distance-based algorithm) works well here.

Step 5: Train the KNN Model

KNN works by looking at the **K nearest neighbors** of a point and letting them “vote” on its class.

```
model = train_knn(X_train, y_train)
```

Example:

If ($K = 5$) and among the 5 nearest users, 3 purchased and 2 did not → the model predicts **Purchased = 1**.

Step 6: Save the Trained Model

We save the trained KNN model using Python's **pickle** for reuse in APIs or apps.

```
save_model(model, '../models/knn_model.pkl')
```

Step 7: Evaluate the Model

We test our model on unseen data and calculate performance metrics.

```
acc, cm = evaluate_model(model, X_test, y_test)

print(f"Accuracy: {acc:.2f}")

print("Confusion Matrix:\n", cm)

sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')

plt.title("Confusion Matrix Heatmap")

plt.xlabel("Predicted")

plt.ylabel("Actual")

plt.show()
```

Interpreting the Confusion Matrix

From your documentation:

```
[[64  4]
 [ 3 29]]
```

Term	Meaning	Count
-----	-----	----
TN (True Negative)	Correctly predicted non-buyers	64
FP (False Positive)	Incorrectly predicted buyers	4
FN (False Negative)	Missed actual buyers	3
TP (True Positive)	Correctly predicted buyers	29

Accuracy: 93%

Precision: 88%

Recall: 91%

F1-Score: 89%

☒ The model performs well — minimal misclassifications, strong generalization.

🧠 Step 8: Visualize Decision Boundaries

Now, let's see **how the model makes its decisions visually**.



```
plot_decision_boundary_train_test(model, X_train, y_train, X_test, y_test)
```

🔗 How It Works (from your visualization doc):

1. **Meshgrid Creation** – a grid of possible (Age, Salary) pairs is generated.


2. **Model Prediction** – KNN predicts each point as "buy" or "not buy."

3. **Color Regions:**

-  **Red:** Model predicts "Not Purchased (0)"
-  **Green:** Model predicts "Purchased (1)"

4. **Data Points:**

- Circles → Training set
- Squares → Test set

 1760813356359

🧠 Graph Insights

- Green and red zones show **decision boundaries** created by KNN.
 - Overlapping regions represent **uncertain areas**.
 - Misclassified points appear in the "wrong" color zone.
 - The model captures complex shapes — a major strength of KNN.
-

☒ Step 9: Conclusion


***KNN** effectively distinguishes buyers from non-buyers based on demographic data.

***Accuracy:** 93% with balanced precision and recall.

***Visualization** reveals how KNN adapts to nonlinear patterns.


***Model saved** as `models/knn_model.pkl` for reuse.

Author Information

 **Author:** Ayoub Majjid

 **Email:** ayoub@majjid.com

 **Portfolio:** majjid.com

 **Repository:***[Insert repo link here]*

 **Company:** INTELLCAP